

1. It is not BCNF; We can first move $F \rightarrow C$ to its own table, i.e. $R_1(A, B, D, E, F)$ and $R_2(F, C)$. Then we decompose R_1 into $R_3(A, B, D, E)$ and $R_4(B, D, F)$ since $BD \rightarrow F$. Therefore the final set of tables is $R_2(F, C), R_3(A, B, D, E), R_4(B, D, F)$.
2. (a) We have $\{C, D, E, F, G, H\}$ that has no dependencies, which makes $2^6 = 64$ closed sets. Then since a closed set with B must also have C, D in it, we have another $2^4 = 16$ sets (since $\{E, F, G, H\}$ makes 16 combinations). Then we can add A to those 16 sets to make 16 more. Therefore we have $64 + 16 + 16 = 96$ closed sets in total.
 (b) Any subset that implies $A \rightarrow B$ must contain either or both $\{A \rightarrow B\}$ and $\{A \rightarrow C, C \rightarrow B\}$. Then we can make 8 combinations that does not logically imply $A \rightarrow B$ with $\{B \rightarrow A, C \rightarrow A, B \rightarrow C\}$, then another 8 from those combinations with $A \rightarrow C$ added (and not $C \rightarrow B$), similarly, we have another 8 with $C \rightarrow B$ added. Therefore the number of subsets that implies $A \rightarrow B$ is $64 - 8 - 8 - 8 = 40$.
3. I: CBAD: guaranteed to empty; Table C can be deleted first since no other table depends on it. Then B can be deleted since although D(z1) depends on B(x) it has ON DELETE SET NULL specified and it is nullable. A can now be deleted since D(z2) is also nullable.
 II: CDAB: not guaranteed to empty; A is deleted before B so B(x) will be set NULL but it is a PRIMARY KEY and cannot be set NULL.
 III: BCDA: guaranteed to empty; B is deleted first and D(z1) is set NULL, C is then deleted since nothing references it. D can then be deleted since A is not dependent on it.

4. (a)

| | INSERT | DELETE | UPDATE |
|-----------|-----------------------------|--------|-------------------------------|
| Executive | YES (insert high pay) | NO | YES (raise pay past 10 times) |
| Employee | YES (insert new lowest pay) | NO | YES (lower lowest pay) |


(b) CHECK (salary <= (SELECT 10*MIN(salary) FROM Employee WHERE division = div_in_charge))

5. SELECT publisher, SUM(price) as a1
FROM Books
WHERE birthYear > 1950
GROUP BY publisher
HAVING SUM(price) > 10000;

6. (a) missing the average rotational latency - r , since the head might not start at the beginning of the file;
Access time = seek time + rotational delay + transfer rate = $10\text{ms} + r + 1/(6000/60) = 10\text{ms} + r + 10\text{ms} = r + 20\text{ms}$.
 (b) $9804; \lfloor \frac{4096 \text{ bytes/block}}{40 \text{ bytes/tuple}} \rfloor = 102 \text{ tuples/block}, \lceil \frac{1000000 \text{ tuples/table}}{102 \text{ tuples/block}} \rceil = 9804 \text{ blocks/table}$.
 (c) Each block is 4096 bytes so each one can fit 255 longs and 256 pointers. So we have $\lceil \frac{1000000}{255} \rceil = 3922$ leaf nodes, $\lceil \frac{3922}{256} \rceil = 15$ level 2 nodes, and $\lceil \frac{15}{256} \rceil = 1$ root node. Therefore we need at least $3922 + 15 + 1 = 3938$ nodes.
 (d) 3 IOs; 2 IOs from root to level 3, 1 IO from level 3 to data.


7. (a)

Height-3 Tree:




(b)

Height-2 Tree:



8. (a)

Insert (3, e)



(b)

Insert (4, f)

